



# Introduction to Learning, Linear regression

---

Iran University of Science and Technology

By: M. S. Tahaei, PhD.

Fall 2024

Courtesy: slides are adopted partly from Dr. Soleymani, Sharif University

# Outline

- Introduction to Learning
- Linear Regression
- Gradient Descent
- Generalized Linear Regression

# A Definition of ML

- ▶ Tom Mitchell (1998): Well-posed learning problem
  - ▶ “A computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E”.
- ▶ Using the observed data to make better decisions
  - ▶ Generalizing from the observed data

# ML Definition: Example

- ▶ Consider an email program that learns how to filter spam according to emails you do or do not mark as spam.
  - ▶ T: Classifying emails as spam or not spam.
  - ▶ E: Watching you label emails as spam or not spam.
  - ▶ P: The number (or fraction) of emails correctly classified as spam/not spam.

# The essence of machine learning

- A pattern exist
- We do not know it mathematically
- We have data on it

# Example: Home Price

- ▶ Housing price prediction




Figure adopted from slides of Andrew Ng,  
Machine Learning course, Stanford.

# Example: Bank loan

- ▶ Applicant form as the input:

age	23 years
gender	male
annual salary	\$30,000
years in residence	1 year
years in job	1 year
current debt	\$15,000
...	...


- ▶ Output: approving or denying the request

# Components of (Supervised) Learning

- Unknown target function:  $f: \mathcal{X} \rightarrow \mathcal{Y}$ 
  - Input space:  $\mathcal{X}$
  - Output space:  $\mathcal{Y}$
- Training data:  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$
- Pick a formula  $g: \mathcal{X} \rightarrow \mathcal{Y}$  that approximates the target function  $f$ 
  - selected from a set of hypotheses  $\mathcal{H}$



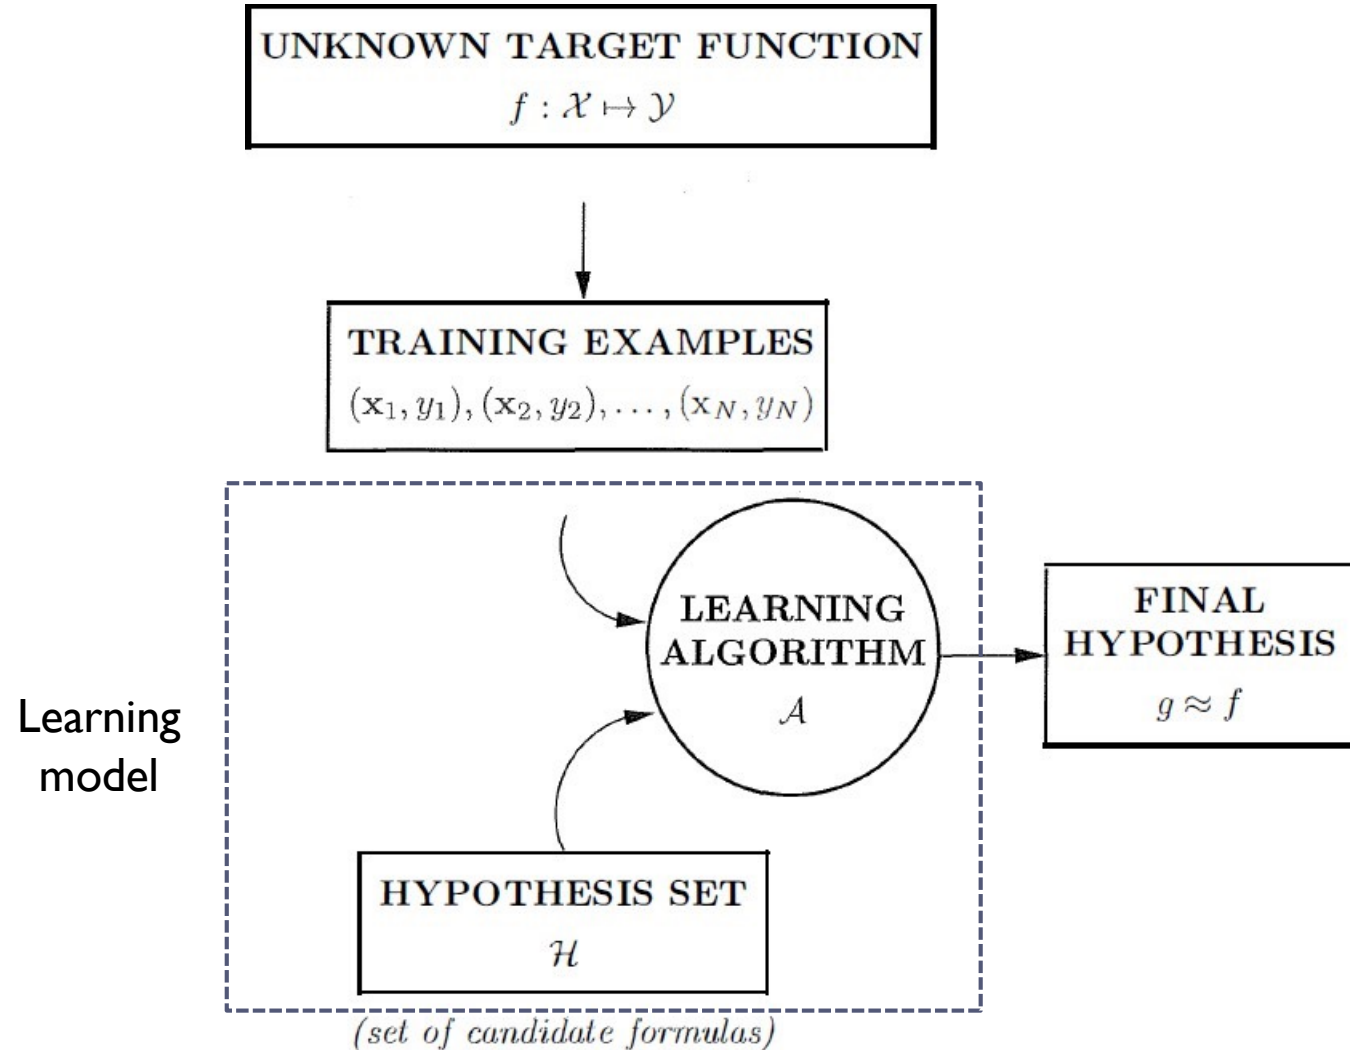
# Training data: Example



Training data

$x_1$	$x_2$	$y$	
0.9	2.3	1	—
3.5	2.6	1	—
2.6	3.3	1	—
2.7	4.1	1	—
1.8	3.9	1	—
6.5	6.8	-1	×
7.2	7.5	-1	×
7.9	8.3	-1	×
6.9	8.3	-1	×
8.8	7.9	-1	×
9.1	6.2	-1	×

# Components of (Supervised) Learning



# Solution Components

- ▶ **Learning model** composed of:
  - ▶ Learning algorithm
  - ▶ Hypothesis set
- ▶ Perceptron example


# Perceptron classifier

- ▶ Input  $\mathbf{x} = [x_1, \dots, x_d]$
- ▶ Classifier:
  - ▶ If  $\sum_{i=1}^d w_i x_i > \text{threshold}$  then output 1
  - ▶ else output -1
- ▶ The linear formula  $g \in \mathcal{H}$  can be written:

$$g(\mathbf{x}) = \text{sign} \left( \sum_{i=1}^d \mathbf{w}_i x_i + \mathbf{w}_0 \right)$$

If we add a coordinate  $x_0 = 1$  to the input:

$$g(\mathbf{x}) = \text{sign} \left( \sum_{i=0}^d \mathbf{w}_i x_i \right) \quad \xrightarrow{\text{Vector form}} \quad g(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$$



# Perceptron learning algorithm: linearly separable data

- ▶ Give the training data  $(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})$
- ▶ **Misclassified** data  $(\mathbf{x}^{(n)}, y^{(n)})$ :  
 $\text{sign}(\mathbf{w}^T \mathbf{x}^{(n)}) \neq y^{(n)}$

Repeat


Pick a **misclassified** data  $(\mathbf{x}^{(n)}, y^{(n)})$  from training data and update  $\mathbf{w}$ :

$$\mathbf{w} = \mathbf{w} + y^{(n)} \mathbf{x}^{(n)}$$





Until all training data points are correctly classified by  $g$

# Perceptron learning algorithm: Example of weight update

$$y = +1$$



$$y = -1$$



# Experience (E) in ML

- ▶ Basic premise of learning:
  - ▶ “Using a set of observations to uncover an underlying process”
- ▶ We have different types of (getting) observations in different types or paradigms of ML methods

# Paradigms of ML

- Supervised learning (regression, classification)
  - predicting a target variable for which we get to see examples.
- Unsupervised learning
  - revealing structure in the observed data
- Reinforcement learning
  - partial (indirect) feedback, no explicit guidance
  - Given rewards for a sequence of moves to learn a policy and utility functions
- Other paradigms: semi-supervised learning, active learning, online learning, etc.



# Supervised Learning: Regression vs. Classification

- Supervised Learning
  - **Regression**: predict a continuous target variable
    - E.g.,  $y \in [0,1]$
  - **Classification**: predict a discrete target variable
    - E.g.,  $y \in \{1,2, \dots, C\}$

# Data in Supervised Learning

- ▶ Data are usually considered as vectors in a  $d$  dimensional space
  - ▶ Now, we make this assumption for illustrative purpose
    - ▶ We will see it is not necessary

Columns:

*Features/attributes/dimensions*

Rows:

*Data/points/instances/examples/samples*

Y column:

*Target/outcome/response/label*

	$x_1$	$x_2$	...	$x_d$	$y$ (Target)
Sample 1					
Sample 2					
...					
Sample n-1					
Sample n					

# Regression: Example

- ▶ Housing price prediction





Figure adopted from slides of Andrew Ng

# Classification: Example

## ► Weight (Cat, Dog)



# Supervised Learning vs. Unsupervised Learning

- ▶ Supervised learning

- ▶ Given: Training set

- ▶ labeled set of  $N$  input-output pairs  $D = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$

- ▶ Goal: learning a mapping from  $\mathbf{x}$  to  $y$

- ▶ Unsupervised learning


- ▶ Given: Training set

- ▶  $\{\mathbf{x}^{(i)}\}_{i=1}^N$

- ▶ Goal: find groups or structures in the data

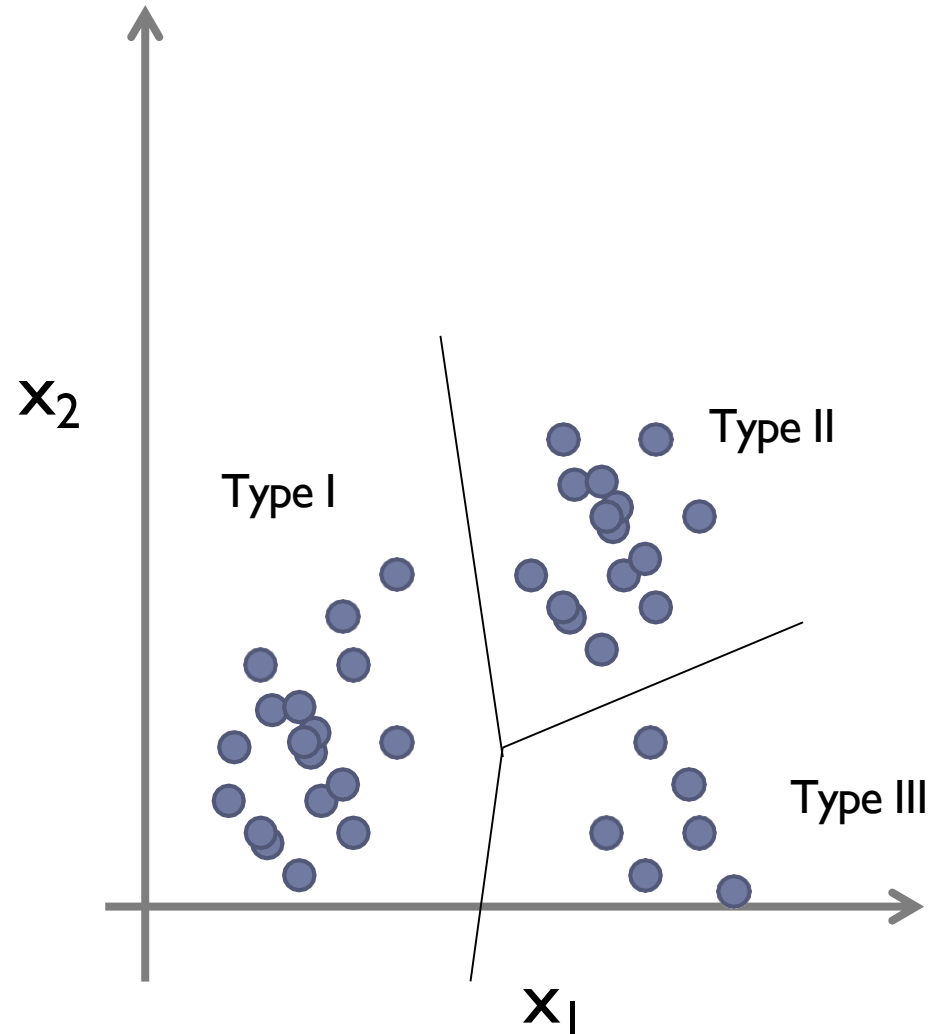
- ▶ Discover the intrinsic structure in the data

# Supervised Learning: Samples



Classification

# Unsupervised Learning: Samples



Clustering

# Sample Data in Unsupervised Learning

## ► Unsupervised Learning:

Columns:

*Features/attributes/dimensions*

Rows:

*Data/points/instances/examples/samples*

	$x_1$	$x_2$	...	$x_d$
Sample 1				
Sample 2				
...				
Sample n-1				
Sample n				



# Unsupervised Learning: Example Applications

- ▶ Clustering docs based on their similarities
  - ▶ Grouping new stories in the Google news site
- ▶ Market segmentation: group customers into different market segments given a database of customer data.
- ▶ Social network analysis


# Reinforcement Learning

Provides only an indication as to whether an action is correct or not




- Data in supervised learning:  
(input, correct output)
- Data in Reinforcement Learning:  
(input, some output, a grade of reward for this output)

# Reinforcement Learning




- Typically, we need to get a sequence of decisions
  - it is usually assumed that reward signals refer to the entire sequence




# Example



$$f = -1$$

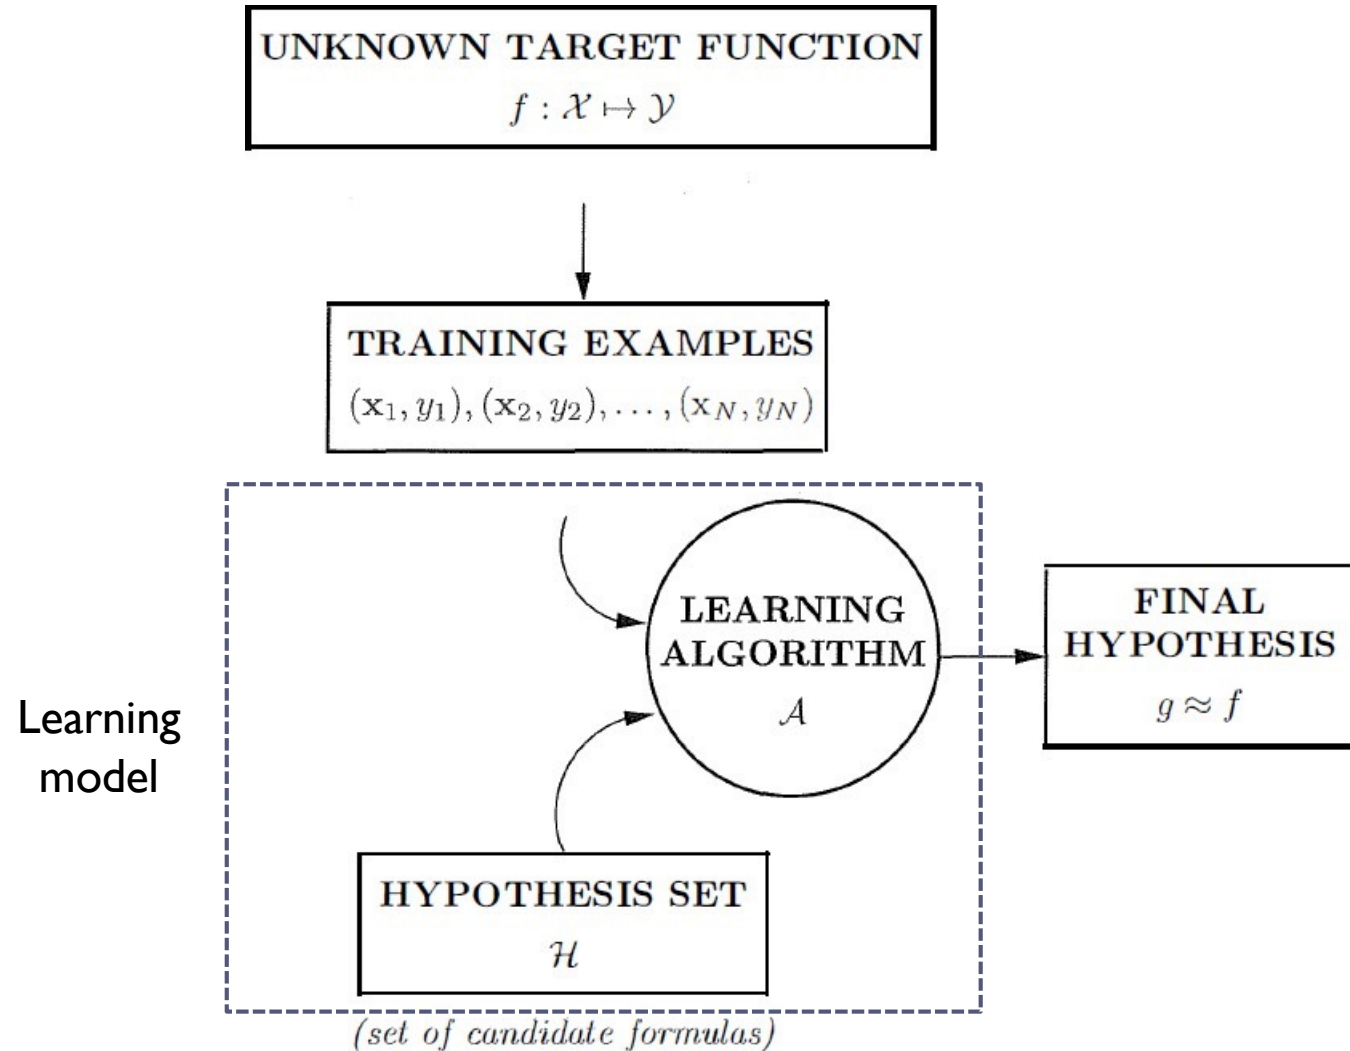


$$f = +1$$



$$f = ?$$

# Components of (Supervised) Learning




# Main Steps of Learning Tasks

- ▶ Selection of hypothesis set (or model specification)
  - ▶ Which class of models (mappings) should we use for our data?
- ▶ Learning: find mapping  $f$  (from hypothesis set) based on the training data
  - ▶ Which notion of error should we use? (loss functions)
  - ▶ Optimization of loss function to find mapping  $f$
- ▶ Evaluation: how well  $f$  generalizes to yet unseen examples
  - ▶ How do we ensure that the error on future data is minimized? (generalization)

# Handwritten Digit Recognition Example

- ▶ **Data: labeled samples**



# Example: Input representation


'raw' input  $\mathbf{x} = (x_0, x_1, x_2, \dots, x_{256})$

linear model:  $(w_0, w_1, w_2, \dots, w_{256})$

**Features:** Extract useful information, e.g.,

intensity and symmetry  $\mathbf{x} = (x_0, x_1, x_2)$

linear model:  $(w_0, w_1, w_2)$






# Example: Illustration of features

$$\mathbf{x} = (x_0, x_1, x_2)$$

$x_1$ : intensity

$x_2$ : symmetry



# Linear regression, Cost Function and Generalization

# Regression problem

- ▶ The goal is to make (real valued) predictions given features
- ▶ Example: predicting house price from 3 attributes

Size ( $m^2$ )	Age (year)	Region	Price ( $10^6T$ )
100	2	5	500
80	25	3	250
...	...	...	...


# Learning problem

- ▶ Selecting a **hypothesis space**
  - ▶ Hypothesis space: a set of mappings from feature vector to target
- ▶ **Learning (estimation)**: optimization of a cost function
  - ▶ Based on the training set  $D = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$  and a cost function we find (an estimate)  $f \in F$  of the target function
- ▶ **Evaluation**: we measure how well  $f$  generalizes to unseen examples

# Linear regression: hypothesis space

- ▶ Univariate

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad f(x; \mathbf{w}) = w_0 + w_1 x$$



- ▶ Multivariate

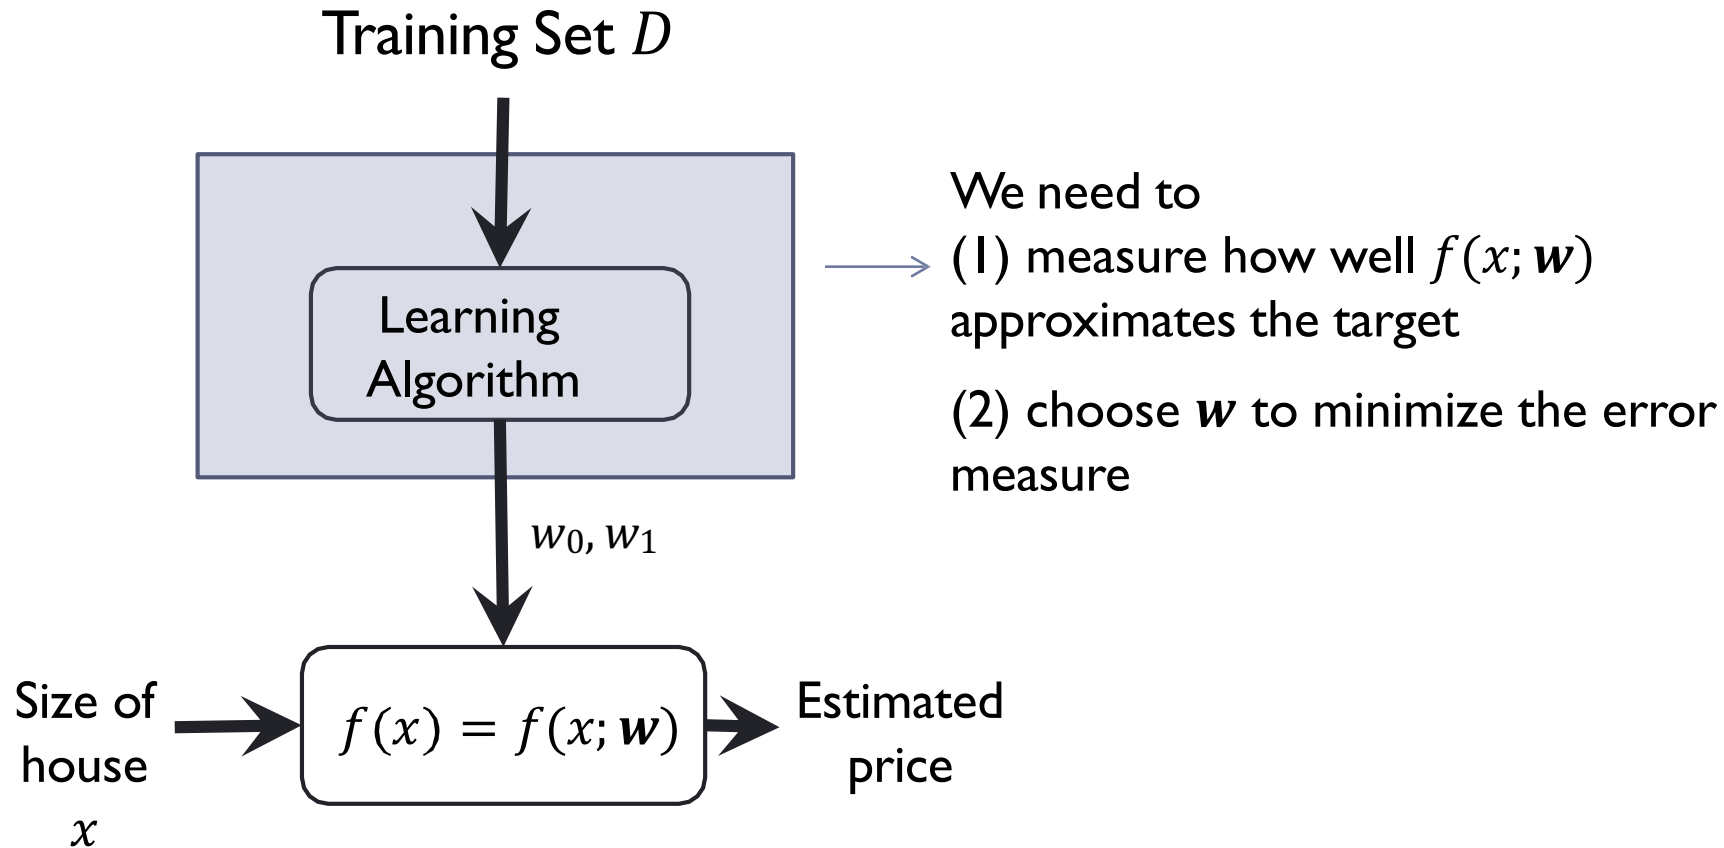
$$f : \mathbb{R}^d \rightarrow \mathbb{R} \quad f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x_1 + \dots w_d x_d$$

$\mathbf{w} = [w_0, w_1, \dots, w_d]^T$  are parameters we need to set.


# Learning algorithm

- ▶ Select how to measure the error (i.e. prediction loss)
- ▶ Find the minimum of the resulting error or cost function

# Learning algorithm




# How to measure the error



Squared error:  $\left(y^{(i)} - f(x^{(i)}; \mathbf{w})\right)^2$



# Linear regression: univariate example



Cost function:

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^n (y^{(i)} - f(x; \mathbf{w}))^2 \\ &= \sum_{i=1}^n (y^{(i)} - w_0 - w_1 x^{(i)})^2 \end{aligned}$$

# Regression: squared loss

- ▶ In the SSE cost function, we used squared error as the prediction loss:

$$Loss(y, \hat{y}) = (y - \hat{y})^2 \quad \hat{y} = f(\mathbf{x}; \mathbf{w})$$

- ▶ Cost function (based on the training set):

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^n Loss(y^{(i)}, f(\mathbf{x}^{(i)}; \mathbf{w})) \\ &= \sum_{i=1}^n \left( y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}) \right)^2 \end{aligned}$$



- ▶ Minimizing sum (or mean) of squared errors is a common approach in curve fitting, neural network, etc.

# Sum of Squares Error (SSE) cost function

$$J(\mathbf{w}) = \sum_{i=1}^n (y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}))^2$$

- ▶  $J(\mathbf{w})$ : sum of the squares of the prediction errors on the training set
- ▶ We want to find the best regression function  $f(\mathbf{x}^{(i)}; \mathbf{w})$ 
  - ▶ equivalently, the best  $\mathbf{w}$
- ▶ Minimize  $J(\mathbf{w})$ 
  - ▶ Find optimal  $f(\mathbf{x}) = f(\mathbf{x}; \mathbf{w}^*)$  where  $\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} J(\mathbf{w})$


# Cost function: univariate example



# Cost function: univariate example


$$f(x; w_0, w_1) = w_0 + w_1 x$$

(for fixed  $w_0, w_1$ , this is a function of  $x$ )




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




# Cost function: univariate example

$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




# Cost function: univariate example

$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




# Cost function: univariate example

$$f(x; w_0, w_1) = w_0 + w_1 x$$



$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )





# Cost function optimization: univariate

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - w_0 - w_1 x^{(i)})^2$$

- ▶ Necessary conditions for the “optimal” parameter values:

$$\frac{\partial J(\mathbf{w})}{\partial w_0} = 0$$

$$\frac{\partial J(\mathbf{w})}{\partial w_1} = 0$$

# Optimality conditions: univariate

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - w_0 - w_1 x^{(i)})^2$$

$$\frac{\partial J(\mathbf{w})}{\partial w_1} = \sum_{i=1}^n (y^{(i)} - w_0 - w_1 x^{(i)})(-x^{(i)}) = 0$$

$$\frac{\partial J(\mathbf{w})}{\partial w_0} = \sum_{i=1}^n (y^{(i)} - w_0 - w_1 x^{(i)})(-1) = 0$$

- ▶ A systems of 2 linear equations

# Cost function: multivariate

- ▶ We have to minimize the empirical squared loss:


$$J(\mathbf{w}) = \sum_{i=1}^n \left( y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}) \right)^2$$

$$f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x_1 + \dots w_d x_d$$

$$\mathbf{w} = [w_0, w_1, \dots, w_d]^T$$

$$\mathbf{w} = \underset{\mathbf{w} \in \mathbb{R}^{d+1}}{\operatorname{argmin}} J(\mathbf{w})$$

# Cost function and optimal linear model



- ▶ Necessary conditions for the “optimal” parameter values:

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \mathbf{0}$$

- ▶ A system of  $d + 1$  linear equations

# Cost function: matrix notation

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^n \left( \mathbf{y}^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}) \right)^2 = \\ &= \sum_{i=1}^n \left( \mathbf{y}^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)} \right)^2 \end{aligned}$$

$$\mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix} \quad \mathbf{X} = \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_d^{(1)} \\ 1 & x_1^{(2)} & \dots & x_d^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(n)} & \dots & x_d^{(n)} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix}$$

$$J(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

# Minimizing cost function

Optimal linear weight vector (for SSE cost function):

$$J(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = 2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \mathbf{0} \Rightarrow \mathbf{X}^T \mathbf{X}\mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

# Minimizing cost function

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y}$$

$$\mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$

$\mathbf{X}^\dagger$  is pseudo inverse of  $\mathbf{X}$

# Another approach for optimizing the sum squared error

- ▶ Iterative approach for solving the following optimization problem:

$$J(\mathbf{w}) = \sum_{i=1}^n (y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}))^2$$



# Gradient descent

- ▶ Cost function:  $J(\mathbf{w})$
- ▶ Optimization problem:  $\mathbf{w} = \underset{\mathbf{w}}{\operatorname{argmin}} J(\mathbf{w})$
- ▶ Steps:
  - ▶ Start from  $\mathbf{w}^0$
  - ▶ Repeat
    - ▶ Update  $\mathbf{w}^t$  to  $\mathbf{w}^{t+1}$  in order to reduce  $J$
    - ▶  $t \leftarrow t + 1$
  - ▶ until we hopefully end up at a minimum

# Gradient descent

- ▶ First-order optimization algorithm to find  $\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} J(\mathbf{w})$ 
  - ▶ Also known as "**steepest descent**"
- ▶ In each step, takes steps proportional to the negative of the gradient vector of the function at the current point  $\mathbf{w}^t$ :

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \gamma_t \nabla J(\mathbf{w}^t)$$

- ▶  $J(\mathbf{w})$  decreases fastest if one goes from  $\mathbf{w}^t$  in the direction of  $-\nabla J(\mathbf{w}^t)$
- ▶ Assumption:  $J(\mathbf{w})$  is defined and differentiable in a neighborhood of a point  $\mathbf{w}^t$


**Gradient ascent** takes steps proportional to (the positive of) the gradient to find a local maximum of the function

# Gradient descent

- ▶ Minimize  $J(\mathbf{w})$

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla_{\mathbf{w}} J(\mathbf{w}^t)$$

Step size  
(Learning rate parameter)




$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \left[ \frac{\partial J(\mathbf{w})}{\partial w_1}, \frac{\partial J(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial J(\mathbf{w})}{\partial w_d} \right]$$

- ▶ If  $\eta$  is small enough, then  $J(\mathbf{w}^{t+1}) \leq J(\mathbf{w}^t)$ .
- ▶  $\eta$  can be allowed to change at every iteration as  $\eta_t$ .

# Gradient descent


- ▶ Local minima problem
- ▶ However, when  $J$  is convex, all local minima are also global minima  $\Rightarrow$  gradient descent can converge to the global solution.

# Problem of gradient descent with non-convex cost functions



# Problem of gradient descent with non-convex cost functions

---



# Gradient descent for SSE cost function

- ▶ Minimize  $J(\mathbf{w})$

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla_{\mathbf{w}} J(\mathbf{w}^t)$$

- ▶  $J(\mathbf{w})$ : Sum of squares error

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n \left( y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}) \right)^2$$

- ▶ Weight update rule for  $f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x}$ :

$$\mathbf{w}^{t+1} = \mathbf{w}^t + \eta \sum_{i=1}^n \left( y^{(i)} - \mathbf{w}^{tT} \mathbf{x}^{(i)} \right) \mathbf{x}^{(i)}$$

# Gradient descent for SSE cost function

- ▶ Weight update rule:  $f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x}$


$$\mathbf{w}^{t+1} = \mathbf{w}^t + \eta \sum_{i=1}^n (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) \mathbf{x}^{(i)}$$

Batch mode: each step  
considers all training data


- ▶  $\eta$ : too small  $\rightarrow$  gradient descent can be slow.
- ▶  $\eta$ : too large  $\rightarrow$  gradient descent can overshoot the minimum. It may fail to converge, or even diverge.




$$f(x; w_0, w_1) = w_0 + w_1 x$$



$J(w_0, w_1)$   
(function of the parameters  $w_0, w_1$ )




$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




This example has been adopted from: Prof. Ng's slides (ML Online Course, Stanford)

$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




This example has been adopted from: Prof. Ng's slides (ML Online Course, Stanford)

$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )




$$f(x; w_0, w_1) = w_0 + w_1 x$$




$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )






$$f(x; w_0, w_1) = w_0 + w_1 x$$



$$J(w_0, w_1)$$

(function of the parameters  $w_0, w_1$ )



# Stochastic gradient descent

- ▶ Example: Linear regression with SSE cost function

$$J^{(i)}(\mathbf{w}) = (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$$

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla_{\mathbf{w}} J^{(i)}(\mathbf{w})$$

$$\mathbf{w}^{t+1} = \mathbf{w}^t + \eta (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) \mathbf{x}^{(i)}$$

Least Mean Squares (LMS)

It is proper for sequential or online learning

# Stochastic gradient descent: online learning

- ▶ Sequential learning is also appropriate for real-time applications
  - ▶ data observations are arriving in a continuous stream
  - ▶ and predictions must be made before seeing all of the data
- ▶ The value of  $\eta$  needs to be chosen with care to ensure that the algorithm converges

# Evaluation and generalization

- ▶ Why minimizing the cost function (based on only training data) while we are interested in the performance on new examples?

$$\min_{\theta} \sum_{i=1}^n \text{Loss} \left( y^{(i)}, f(\mathbf{x}^{(i)}; \theta) \right) \longrightarrow \text{Empirical loss}$$

- ▶ **Evaluation:** After training, we need to measure how well the learned prediction function can predicts the target for unseen examples

# Training and test performance





- ▶ Assumption: training and test examples are drawn independently at random from the same but unknown distribution.
  - ▶ Each training/test example  $(x, y)$  is a sample from joint probability distribution  $P(x, y)$ , i.e.,  $(x, y) \sim P$

$$\textbf{Empirical (training) loss} = \frac{1}{n} \sum_{i=1}^n \text{Loss} \left( y^{(i)}, f(x^{(i)}; \theta) \right)$$

$$\textbf{Expected (test) loss} = E_{x,y} \{ \text{Loss}(y, f(x; \theta)) \}$$


- ▶ We minimize empirical loss (on the training data) and expect to also find an acceptable expected loss
  - ▶ Empirical loss as a proxy for the performance over the whole distribution.

# Linear regression: number of training data



# Linear regression: generalization

- ▶ By increasing the number of training examples, will solution be better?
- ▶ Why the mean squared error does not decrease more after reaching a level?



# Linear regression: types of errors

- ▶ Structural error: the error introduced by the limited function class (infinite training data):

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} E_{x,y} [(y - \mathbf{w}^T \mathbf{x})^2]$$

$$\text{Structural error: } E_{x,y} \left[ (y - \mathbf{w}^{*T} \mathbf{x})^2 \right]$$

where  $\mathbf{w}^* = (w_0^*, \dots, w_d^*)$  are the optimal linear regression parameters (infinite training data)



# Linear regression: types of errors

- ▶ Approximation error measures how close we can get to the optimal linear predictions with limited training data:

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} E_{\mathbf{x},y}[(y - \mathbf{w}^T \mathbf{x})^2]$$

$$\mathbf{w} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n \left( y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)} \right)^2$$

$$\text{Approximation error: } E_{\mathbf{x}} \left[ \left( \mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x} \right)^2 \right]$$

Where  $\mathbf{w}$  are the parameter estimates based on a small training set (so themselves are random variables).

# Linear regression: error decomposition

- ▶ The expected error can decompose into the sum of structural and approximation errors

$$\begin{aligned} E_{x,y}[(y - \mathbf{w}^T \mathbf{x})^2] \\ = E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x})^2] + E_x[(\mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})^2] \end{aligned}$$

- ▶ Derivation

$$\begin{aligned} E_{x,y}[(y - \mathbf{w}^T \mathbf{x})^2] &= E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x} + \mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})^2] \\ &= E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x})^2] + E_x[(\mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})^2] \\ &\quad + 2E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x})(\mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})] \end{aligned}$$

# Linear regression: error decomposition

- ▶ The expected error can decompose into the sum of structural and approximation errors

$$\begin{aligned} E_{x,y}[(y - \mathbf{w}^T \mathbf{x})^2] \\ = E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x})^2] + E_x[(\mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})^2] \end{aligned}$$

- ▶ Derivation

$$\begin{aligned} E_{x,y}[(y - \mathbf{w}^T \mathbf{x})^2] &= E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x} + \mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})^2] \\ &= E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x})^2] + E_x[(\mathbf{w}^{*T} \mathbf{x} - \mathbf{w}^T \mathbf{x})^2] \\ &\quad + 0 \end{aligned}$$

Note: Optimality condition for  $\mathbf{w}^*$  give us  $E_{x,y}[(y - \mathbf{w}^{*T} \mathbf{x})\mathbf{x}] = 0$   
since  $\nabla_{\mathbf{w}} E_{x,y}[(y - \mathbf{w}^T \mathbf{x})^2]_{\mathbf{w}^*} = 0$

# Recall: Linear regression (squared loss)

- ▶ Linear regression functions

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad f(x; \mathbf{w}) = w_0 + w_1 x$$

$$f : \mathbb{R}^d \rightarrow \mathbb{R} \quad f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x_1 + \dots w_d x_d$$

$\mathbf{w} = [w_0, w_1, \dots, w_d]^T$  are the parameters we need to set.

- ▶ Minimizing the squared loss for linear regression

$$J(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$

- ▶ We obtain  $\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

# Beyond linear regression

- ▶ How to extend the linear regression to non-linear functions?
  - ▶ Transform the data using basis functions
  - ▶ Learn a linear regression on the new feature vectors (obtained
    - by basis functions)

# Beyond linear regression





- ▶  $m^{th}$  order polynomial regression (univariate  $f : \mathbb{R} \rightarrow \mathbb{R}$ )

$$f(x; \mathbf{w}) = w_0 + w_1x + \dots + w_{m-1}x^{m-1} + w_mx^m$$

- ▶ Solution:  $\mathbf{w} = (\mathbf{X}'^T \mathbf{X}')^{-1} \mathbf{X}'^T \mathbf{y}$

$$\mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \quad \mathbf{X}' = \begin{bmatrix} 1 & x^{(1)1} & x^{(1)2} & \dots & x^{(1)m} \\ 1 & x^{(2)1} & x^{(2)2} & \dots & x^{(2)m} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x^{(n)1} & x^{(n)2} & \dots & x^{(n)m} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_m \end{bmatrix}$$

# Polynomial regression: example



# Generalized linear

- ▶ Linear combination of fixed non-linear function of the input vector

$$f(\mathbf{x}; \mathbf{w}) = w_0 + w_1\phi_1(\mathbf{x}) + \dots w_m\phi_m(\mathbf{x})$$

$\{\phi_1(\mathbf{x}), \dots, \phi_m(\mathbf{x})\}$ : set of basis functions (or features)

$$\phi_i(\mathbf{x}): \mathbb{R}^d \rightarrow \mathbb{R}$$



# Basis functions: examples

- ▶ Linear


If  $m = d$ ,  $\phi_i(\mathbf{x}) = x_i$ ,  $i = 1, \dots, d$ , then

$$f(\mathbf{x}; \mathbf{w}) = w_0 + w_1x_1 + \dots + w_dx_d$$

- ▶ Polynomial (univariate)

If  $\phi_i(x) = x^i$ ,  $i = 1, \dots, m$ , then

$$f(x; \mathbf{w}) = w_0 + w_1x + \dots + w_{m-1}x^{m-1} + w_mx^m$$



# Generalized linear: optimization

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^n \left( y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}) \right)^2 \\ &= \sum_{i=1}^n \left( y^{(i)} - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}^{(i)}) \right)^2 \end{aligned}$$

$$\mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix} \quad \boldsymbol{\Phi} = \begin{bmatrix} 1 & \phi_1(\mathbf{x}^{(1)}) & \cdots & \phi_m(\mathbf{x}^{(1)}) \\ 1 & \phi_1(\mathbf{x}^{(2)}) & \cdots & \phi_m(\mathbf{x}^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \phi_1(\mathbf{x}^{(n)}) & \cdots & \phi_m(\mathbf{x}^{(n)}) \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_m \end{bmatrix}$$

$$\mathbf{w} = (\boldsymbol{\Phi}^T \boldsymbol{\Phi})^{-1} \boldsymbol{\Phi}^T \mathbf{y}$$

# Resource

- 1 C. M. Bishop, *Pattern Recognition and Machine Learning*.
- 2 Y. S. Abu-Mostafa, “Machine learning.” California Institute of Technology, 2012.