



# Matrix Factorization & Collaborative Filtering

Iran University of Science and Technology

M. S. Tahaei PhD.

Fall 2024

Machine Learning, Carnegie Mellon University

# Outline

- **Recommender Systems**
  - Content Filtering
  - Collaborative Filtering (CF)
    - CF: Neighborhood Methods
    - CF: Latent Factor Methods
- **Matrix Factorization**
  - Optimization problem – SGD
  - SGD with Regularization
  - Alternating Least Squares
  - User/item bias terms (matrix trick)

# Recommender Systems

## A Common Challenge:

- Assume you're a company selling **items** of some sort: movies, songs, products, etc.
- Company collects millions of **ratings** from **users** of their **items**
- To maximize profit / user happiness, you want to **recommend** items that users are likely to want

# Recommender Systems

The screenshot shows the Amazon homepage with a Cyber Monday deal banner. A user profile for "Matt" is displayed, along with a "Sign In" button. The main content area features a "Recommended for you, Matt" section with four categories of products:

- Buy It Again in Grocery**: Includes items like pancake mix, maple syrup, Jif peanut butter, and coffee.
- Buy It Again in Pets**: Includes pet supplies like dog treats, cat litter, and cat food.
- Buy It Again in Baby Products**: Includes pacifiers, diapers, and baby toys.
- Engineering Books**: A category for books on probabilistic graphical models.

# Recommender Systems

- **Setup:**

- **Items:**

- movies, songs, products, etc.  
(often many thousands)

- **Users:**

- watchers, listeners, purchasers, etc.  
(often many millions)

- **Feedback:**

- 5-star ratings, not-clicking ‘next’,  
purchases, etc.

- **Key Assumptions:**

- Can represent ratings numerically  
as a user/item matrix
  - Users only rate a small number of  
items (the matrix is sparse)

	Doctor Strange	Star Trek: Beyond	Zootopia
Alice	1		5
Bob	3	4	
Charlie	3	5	2

# Types of Recommender Systems

## Content Filtering

- **Example:** Pandora.com music recommendations (Music Genome Project)
- **Con:** Assumes access to **side information** about items (e.g. properties of a song)
- **Pro:** Got a **new item** to add? No problem, just be sure to include the side information

## Collaborative Filtering

- **Example:** Netflix movie recommendations
- **Pro:** Does not assume access to **side information** about items (e.g. does not need to know about movie genres)
- **Con:** Does not work on **new items** that have no ratings

# Types of Recommender Systems

## Content Filtering

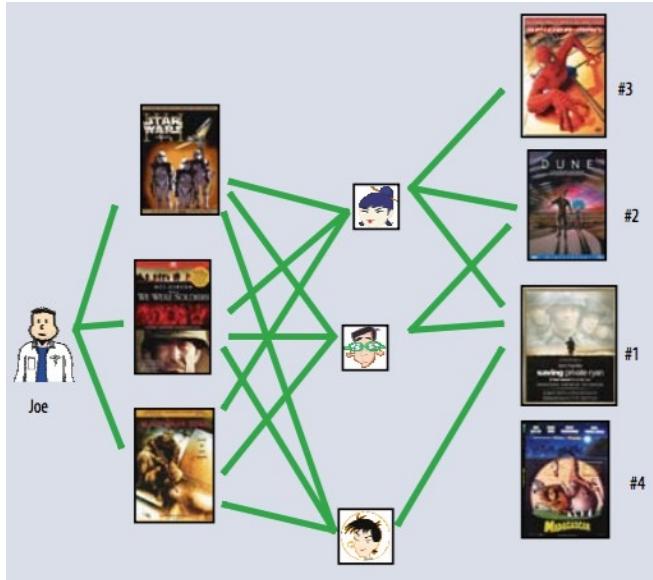
- *Example:* **Pandora.com** music recommendations (Music Genome Project)
- **Con:** Assumes access to **side information** about items (e.g. properties of a song)
- **Pro:** Got a **new item** to add? No problem, just be sure to include the side information

# Collaborative Filtering

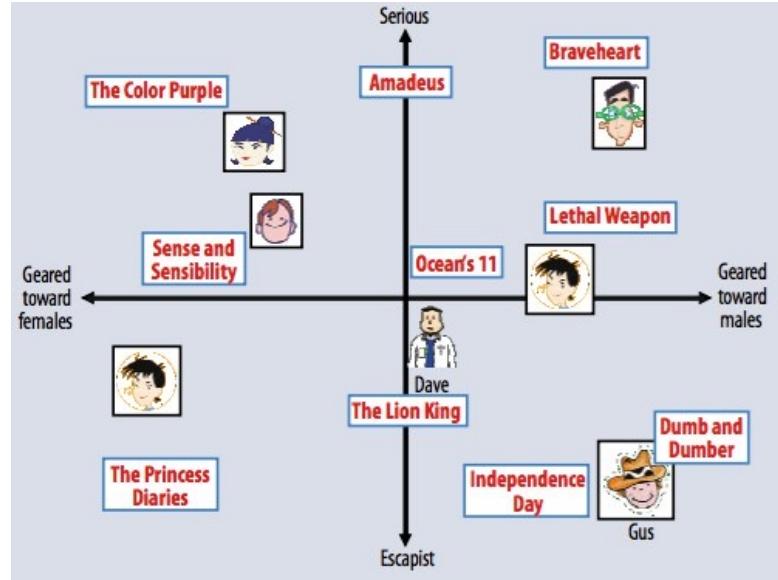
- **Everyday Examples of Collaborative Filtering...**
  - Bestseller lists
  - Top 40 music lists
  - The “recent returns” shelf at the library
  - Unmarked but well-used paths thru the woods
  - The printer room at work
  - “Read any good books lately?”
  - ...
- **Common insight:** personal tastes are correlated
  - If Alice and Bob both like X and Alice likes Y then Bob is more likely to like Y
  - especially (perhaps) if Bob knows Alice

# Two Types of Collaborative Filtering

## 1. Neighborhood Methods



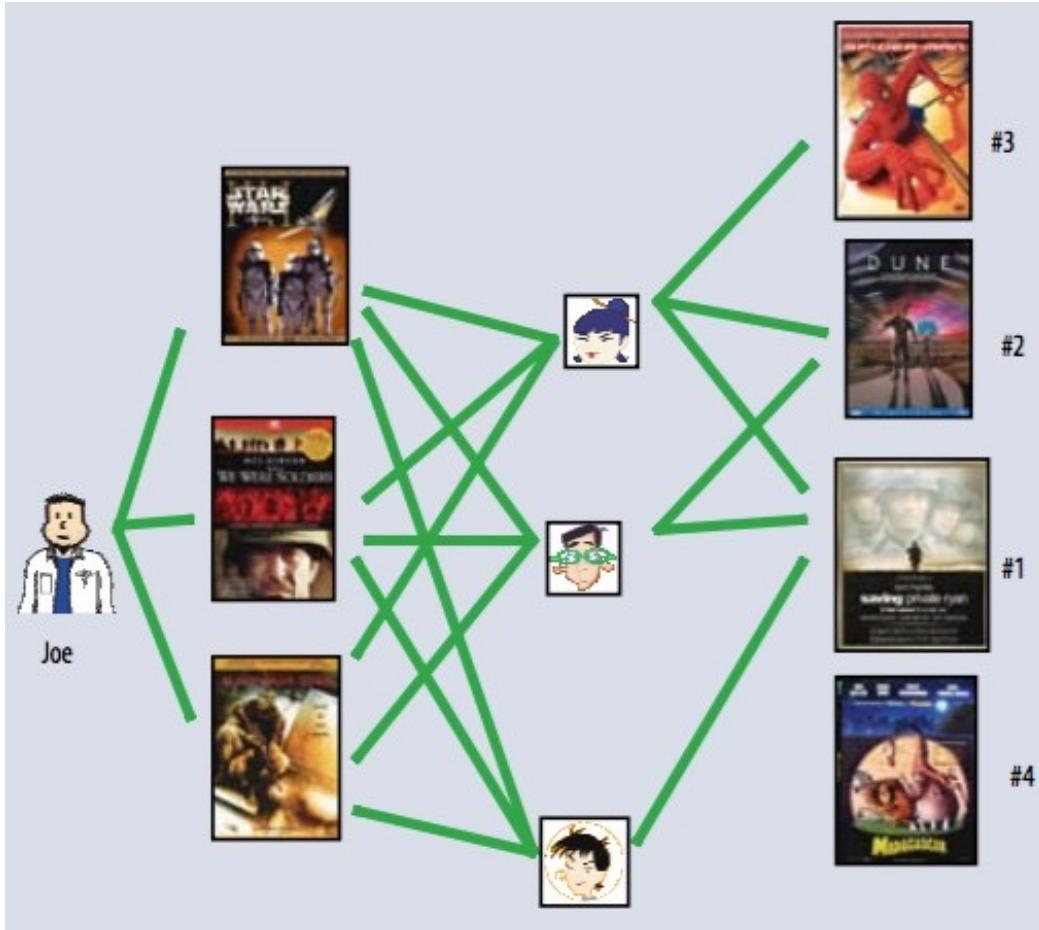
## 2. Latent Factor Methods



Figures from Koren et al. (2009)

# Two Types of Collaborative Filtering

## 1. Neighborhood Methods

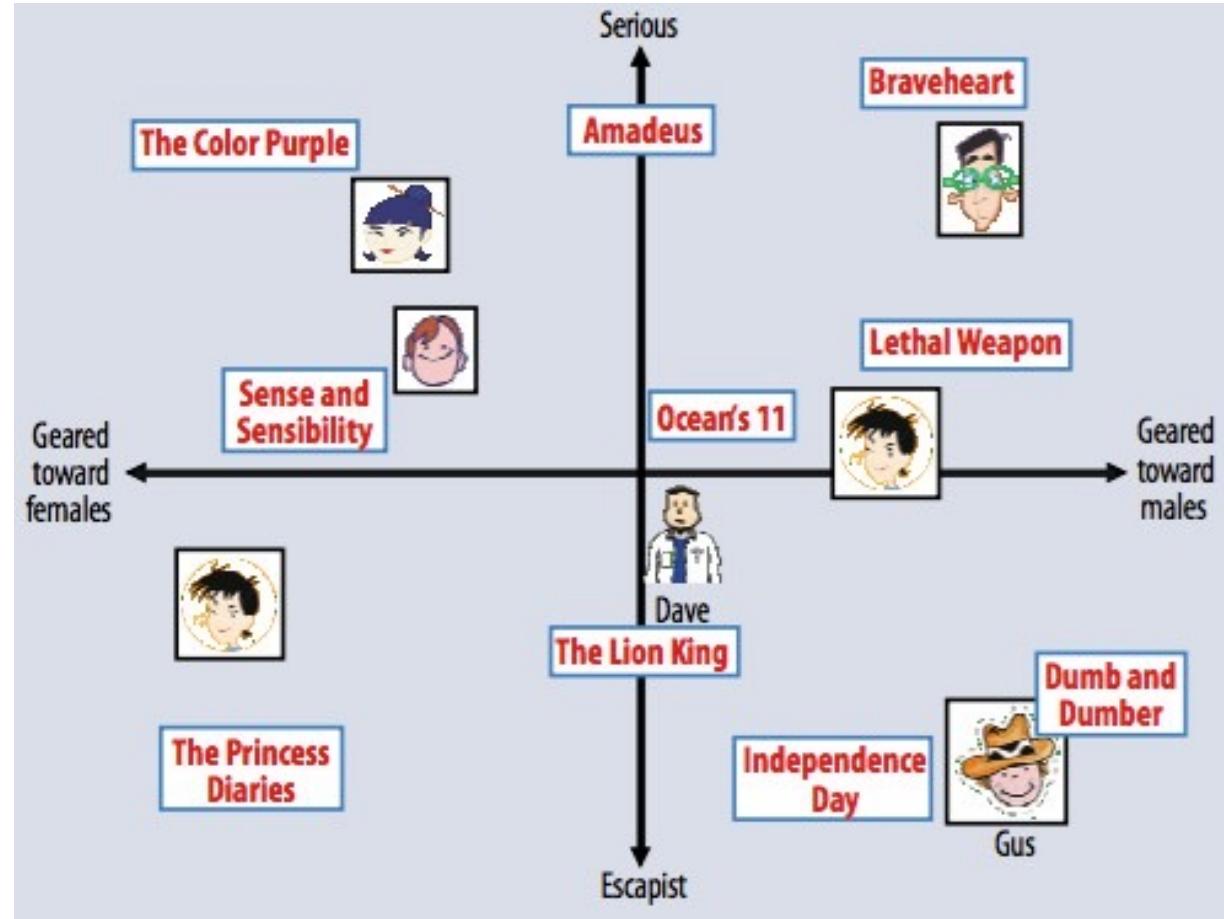


- In the figure, assume that a green line indicates the movie was **watched**
- **Algorithm:**
  1. **Find neighbors** based on similarity of movie preferences
  2. **Recommend** movies that those neighbors watched

# Two Types of Collaborative Filtering

- Assume that both movies and users live in some **low-dimensional space** describing their properties
- **Recommend** a movie based on its **proximity** to the user in the latent space

## 2. Latent Factor Methods



# Matrix Factorization

- Many different ways of factorizing a matrix
- We'll consider three:
  1. Unconstrained Matrix Factorization
  2. Singular Value Decomposition
  3. Non-negative Matrix Factorization
- MF is just another example of a **common recipe**:
  1. define a model
  2. define an objective function
  3. optimize with SGD

# Matrix Factorization

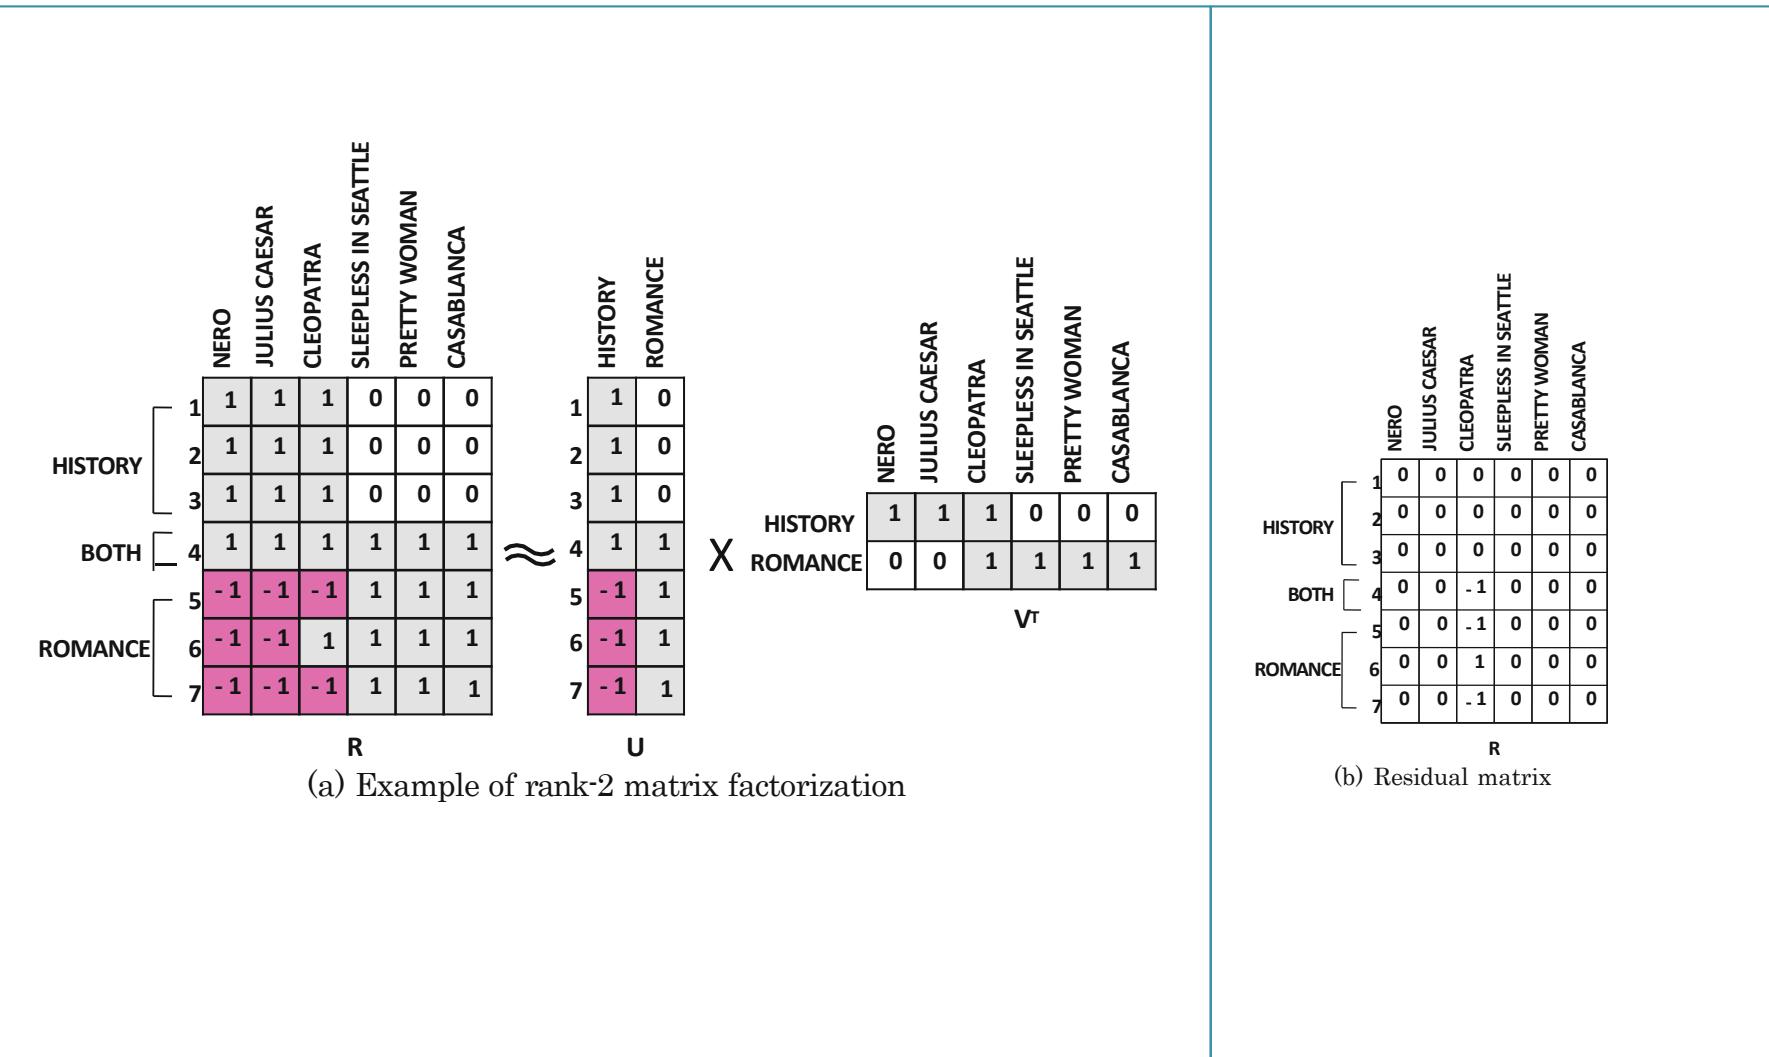
## **Question:**

Applied to the Netflix Prize problem, which of the following methods *always* requires side information about the users and movies?

## **Select all that apply**

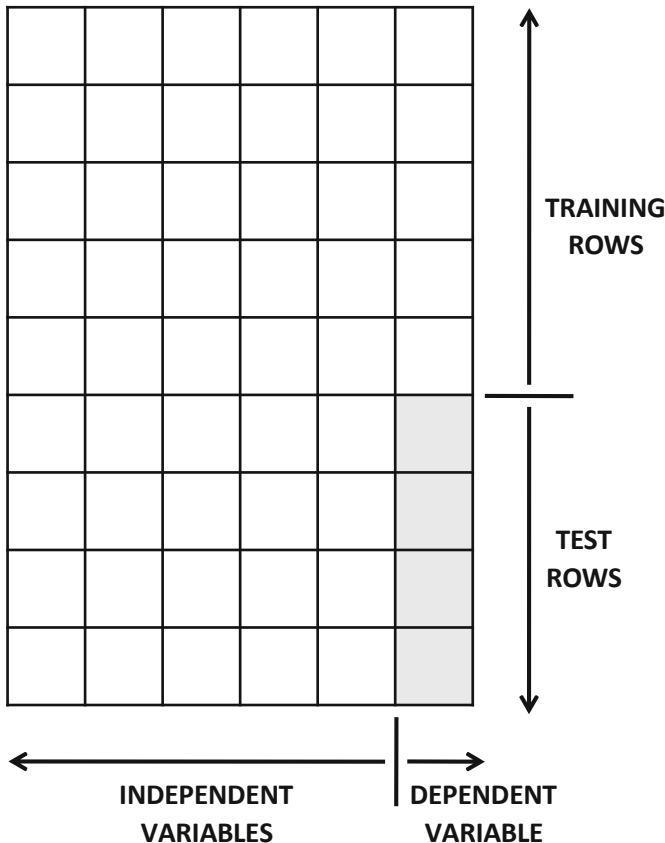
- 1.K-Means
- 2.collaborative filtering
- 3.latent factor methods
- 4.ensemble methods
- 5.content filtering
- 6.neighborhood methods
- 7.recommender systems

# Example: MF for Netflix Problem

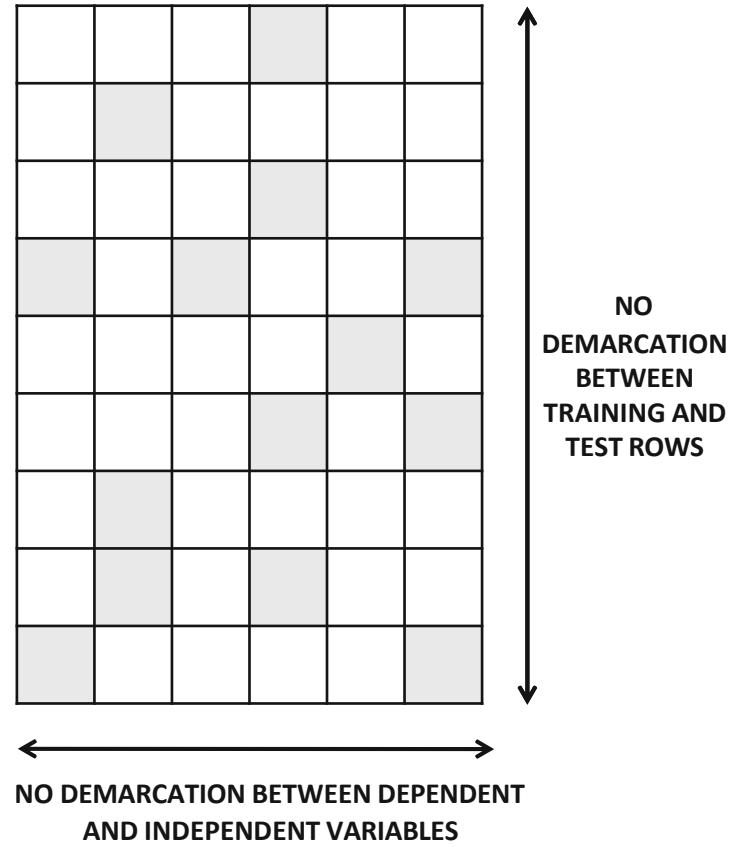


# Regression vs. Collaborative Filtering

**Regression**



**Collaborative Filtering**



# Unconstrained Matrix Factorization

- Optimization problem
- SGD
- SGD with Regularization
- Alternating Least Squares
- User/item bias terms (matrix trick)

# Unconstrained Matrix Factorization

## In-Class Exercise

Derive a block coordinate descent algorithm for the Unconstrained Matrix Factorization problem.

- User vectors:  
 $\mathbf{w}_u \in \mathbb{R}^r$
- Item vectors:  
 $\mathbf{h}_i \in \mathbb{R}^r$
- Rating prediction:  
 $v_{ui} = \mathbf{w}_u^T \mathbf{h}_i$
- Set of non-missing entries  
 $\mathcal{Z} = \{(u, i) : v_{ui} \text{ is observed}\}$
- Objective:  
$$\operatorname{argmin}_{\mathbf{w}, \mathbf{h}} \sum_{(u, i) \in \mathcal{Z}} (v_{ui} - \mathbf{w}_u^T \mathbf{h}_i)^2$$

# Matrix Factorization (with vectors)

- User vectors:

$$\mathbf{w}_u \in \mathbb{R}^r$$

- Item vectors:

$$\mathbf{h}_i \in \mathbb{R}^r$$

- Rating prediction:

$$v_{ui} = \mathbf{w}_u^T \mathbf{h}_i$$

- Set of non-missing entries:

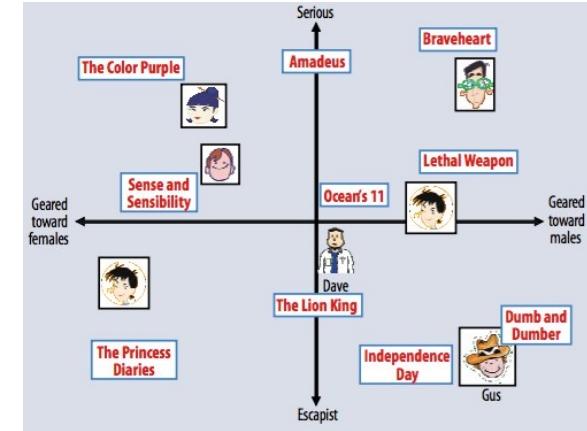
$$\mathcal{Z} = \{(u, i) : v_{ui} \text{ is observed}\}$$

- Objective:

$$\operatorname{argmin}_{\mathbf{w}, \mathbf{h}} \sum_{(u, i) \in \mathcal{Z}} (v_{ui} - \mathbf{w}_u^T \mathbf{h}_i)^2$$

- Regularized Objective:

$$\begin{aligned} \operatorname{argmin}_{\mathbf{w}, \mathbf{h}} \sum_{(u, i) \in \mathcal{Z}} (v_{ui} - \mathbf{w}_u^T \mathbf{h}_i)^2 \\ + \lambda \left( \sum_i \|\mathbf{w}_i\|^2 + \sum_u \|\mathbf{h}_u\|^2 \right) \end{aligned}$$



Figures from Koren et al. (2009)

- SGD update for random (u,i):

$$e_{ui} \leftarrow v_{ui} - \mathbf{w}_u^T \mathbf{h}_i$$

$$\mathbf{w}_u \leftarrow \mathbf{w}_u + \gamma(e_{ui} \mathbf{h}_i - \lambda \mathbf{w}_u)$$

$$\mathbf{h}_i \leftarrow \mathbf{h}_i + \gamma(e_{ui} \mathbf{w}_u - \lambda \mathbf{h}_i)$$

# Matrix Factorization (with matrices)

- User vectors:

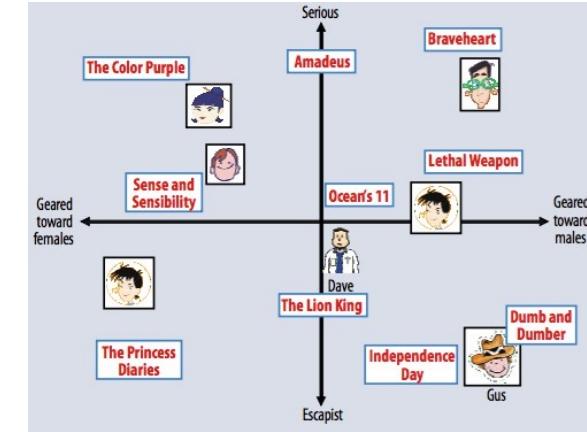
$$(W_u)^T \in \mathbb{R}^r$$

- Item vectors:

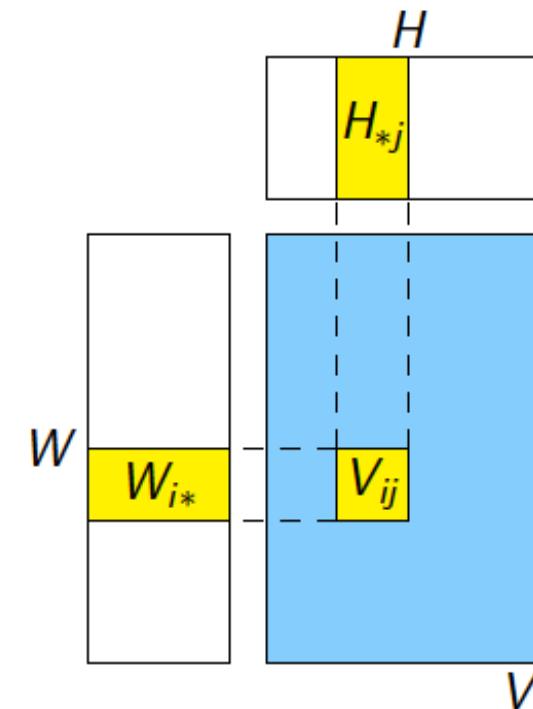
$$H_i \in \mathbb{R}^r$$

- Rating prediction:

$$\begin{aligned} V_{ui} &= W_u H_i \\ &= [WH]_{ui} \end{aligned}$$



Figures from Koren et al. (2009)



# Matrix Factorization

## (with matrices)

- SGD

require that the loss can be written as

$$L = \sum_{(i,j) \in Z} l(V_{ij}, W_{i*}, H_{*j})$$

---

### Algorithm 1 SGD for Matrix Factorization

---

**Require:** A training set  $Z$ , initial values  $W_0$  and  $H_0$

**while** not converged **do** {step}

    Select a training point  $(i, j) \in Z$  uniformly at random.

$$W'_{i*} \leftarrow W_{i*} - \epsilon_n N \frac{\partial}{\partial W_{i*}} l(V_{ij}, W_{i*}, H_{*j})$$

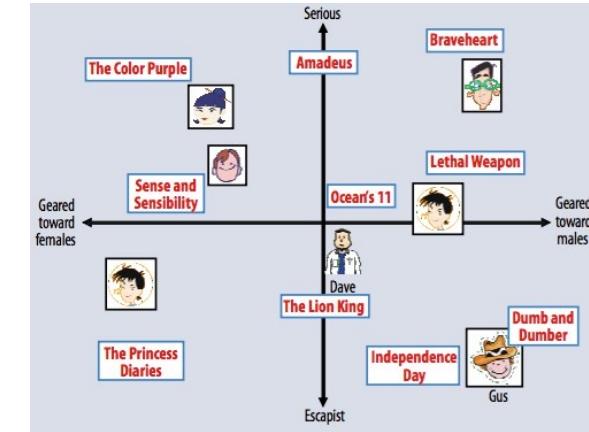
$$H_{*j} \leftarrow H_{*j} - \epsilon_n N \frac{\partial}{\partial H_{*j}} l(V_{ij}, W_{i*}, H_{*j})$$

$$W_{i*} \leftarrow W'_{i*}$$

**end while**

---

step size



Figures from Koren et al. (2009)

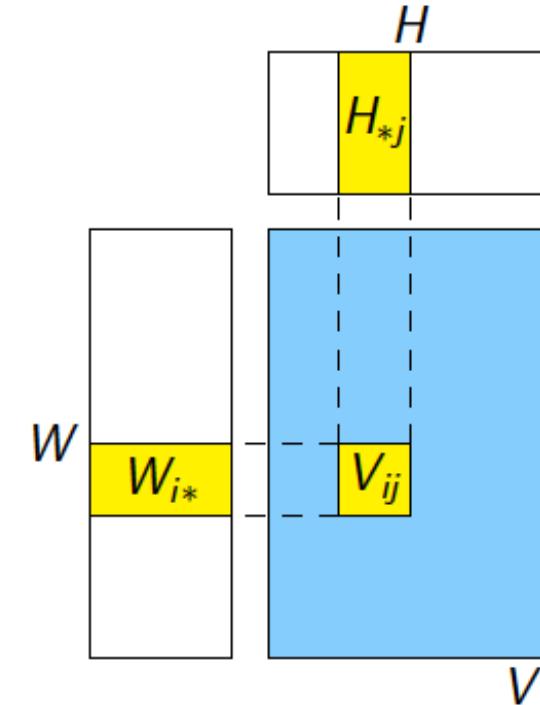


Figure from Gemulla et al. (2011)

# Singular Value Decomposition for Collaborative Filtering

For any arbitrary matrix  $\mathbf{A}$ , SVD gives a decomposition:

$$\mathbf{A} = \mathbf{U}\Lambda\mathbf{V}^T$$

where  $\Lambda$  is a diagonal matrix, and  $\mathbf{U}$  and  $\mathbf{V}$  are orthogonal matrices.

Suppose we have the SVD of our ratings matrix

$$R = Q\Sigma P^T,$$

but then we truncate each of  $Q$ ,  $\Sigma$ , and  $P$  s.t.  $Q$  and  $P$  have only  $k$  columns and  $\Sigma$  is  $k \times k$ :

$$R \approx Q_k \Sigma_k P_k^T$$

For collaborative filtering, let:

$$U \triangleq Q_k \Sigma_k$$

$$V \triangleq P_k$$

$$\Rightarrow U, V = \underset{U, V}{\operatorname{argmin}} \frac{1}{2} \|R - UV^T\|_2^2$$

s.t. columns of  $U$  are mutually orthogonal

s.t. columns of  $V$  are mutually orthogonal

# Non-negative Matrix Factorization

**Constrained Optimization Problem:**

$$U, V = \operatorname{argmin}_{U, V} \frac{1}{2} \|R - UV^T\|_2^2$$

$$\text{s.t. } U_{ij} \geq 0$$

$$\text{s.t. } V_{ij} \geq 0$$

**Multiplicative Updates:** simple iterative algorithm for solving just involves multiplying a few entries together

# SVD vs MF

**Theorem:** If  $R$  fully observed and no regularization, the optimal  $UV^T$  from SVD equals the optimal  $UV^T$  from Unconstrained MF

# Alternating Least Squares (ALS) for UMF:

## Alternating Least Squares (ALS)

(1)

### Cost function $J$

Given  $x^{(1)}, \dots, x^{(n_m)}$ , estimate  $\theta^{(1)}, \dots, \theta^{(n_u)}$ :

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

### Gradient Update

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left( \sum_{i : R^{(i,j)} \neq \text{NA}} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)} \right)$$

(2)

### Cost function $J$

Given  $\theta^{(1)}, \dots, \theta^{(n_u)}$ , learn  $x^{(1)}, \dots, x^{(n_m)}$ :

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

### Gradient Update

$$x_k^{(i)} := x_k^{(i)} - \alpha \left( \sum_{j : R^{(i,j)} \neq \text{NA}} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) \theta_k^{(j)} + \lambda x_k^{(i)} \right)$$

Iterate

# Matrix Factorization

## Comparison of Optimization Algorithms

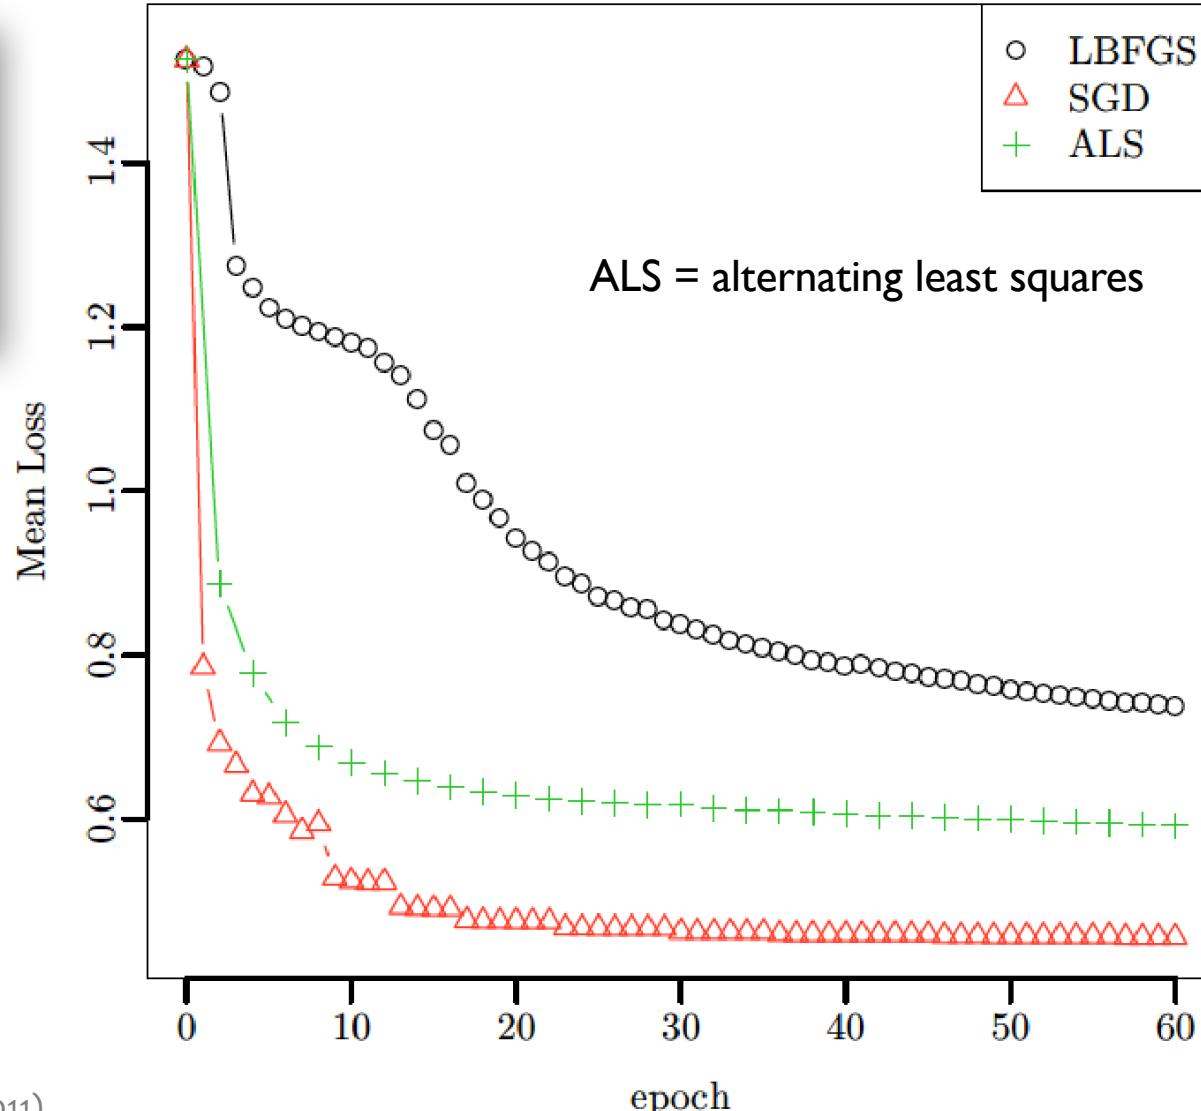


Figure from Gemulla et al. (2011)

# Implicit Feedback Datasets

- What information does a five-star rating contain?
- Implicit Feedback Datasets:
  - In many settings, users don't have a way of expressing *dislike* for an item (e.g. can't provide negative ratings)
  - The only mechanism for feedback is to “like” something
- Examples:
  - Facebook has a “Like” button, but no “Dislike” button
  - Google's “+1” button
  - Pinterest pins
  - Purchasing an item on Amazon indicates a preference for it, but there are many reasons you might not purchase an item (besides dislike)
  - Search engines collect click data but don't have a clear mechanism for observing dislike of a webpage

# Summary

- Recommender systems solve many **real-world** (\*large-scale) **problems**
- Collaborative filtering by Matrix Factorization (MF) is an **efficient** and **effective** approach
- MF is just another example of a **common recipe**:
  1. define a model
  2. define an objective function
  3. optimize with SGD